

9 EXTENDED APPENDIX

9.1 REDS performance on DNN and CNN architectures

In addition to the results on DS-CNN reported in the main paper, we show in Table 5 and Table 6 REDS performance on DNN and CNN architectures (with full fine-tuning) and compare to training model of each capacity from scratch and training REDS from scratch. Despite full fine-tuning, the results for S architecture show superior performance of the BU heuristic over TD.

MACs	Small (S) - Accuracy 83.82				Large (L) - Accuracy 86.87			
	Scratch	Knapsack BU	Knapsack TD	REDS training	Scratch	Knapsack BU	Knapsack TD	REDS training
100%	84.30 \pm 0.11	83.52 \pm 0.07	82.80 \pm 0.16	82.13 \pm 0.20	86.54 \pm 0.24	86.46 \pm 0.34	86.25 \pm 0.19	85.06 \pm 0.19
75%	83.77 \pm 0.23	82.29 \pm 0.35	81.88 \pm 0.30	81.23 \pm 0.21	85.96 \pm 0.13	86.38 \pm 0.65	86.09 \pm 0.03	84.93 \pm 0.20
50%	80.91 \pm 0.11	78.36 \pm 1.40	78.59 \pm 0.24	77.05 \pm 0.34	85.24 \pm 0.35	85.62 \pm 0.24	85.58 \pm 0.35	84.08 \pm 0.22
25%	69.77 \pm 0.67	64.42 \pm 1.99	61.43 \pm 3.35	63.69 \pm 3.26	84.22 \pm 0.13	82.61 \pm 0.61	83.00 \pm 0.62	82.03 \pm 0.59

Table 5: Test set accuracy [%] of training S and L fully-connected (DNN) architectures taken from [53]: training a network of each size from scratch ("Scratch"), conversion from a pre-trained network using two knapsack versions ("Knapsack BU" and "Knapsack TD"), and training REDS structure from scratch ("REDS training"). Reported results from three independent runs. The accuracy of each 100 % network reported in [53] is listed in the header row.

MACs	Small (S) - Accuracy 92.24				Large (L) - Accuracy 93.24			
	Scratch	Knapsack BU	Knapsack TD	REDS training	Scratch	Knapsack BU	Knapsack TD	REDS training
100%	91.10 \pm 0.23	91.60 \pm 0.39	91.20 \pm 0.35	88.89 \pm 0.26	92.94 \pm 0.20	92.83 \pm 0.26	92.97 \pm 0.15	90.97 \pm 0.21
75%	90.40 \pm 0.27	90.63 \pm 0.19	90.20 \pm 0.13	87.64 \pm 0.46	92.74 \pm 0.12	92.74 \pm 0.23	92.54 \pm 0.07	90.67 \pm 0.09
50%	89.07 \pm 0.24	88.39 \pm 0.31	88.39 \pm 0.52	85.99 \pm 0.19	92.44 \pm 0.30	91.95 \pm 0.22	91.93 \pm 0.28	90.36 \pm 0.30
25%	82.57 \pm 0.41	79.52 \pm 0.67	79.28 \pm 0.54	79.25 \pm 0.40	90.98 \pm 0.60	90.24 \pm 0.35	90.31 \pm 0.03	88.25 \pm 0.66

Table 6: The same as in Table 5 for S and L convolutional architectures (CNN) from [53].

Fig. 10 shows the impact of the architecture on REDS structure found by the knapsack BU solver. We present the results for DNN S, L and CNN S, L from left to right, respectively.




Figure 10: Analysis of the subnetwork architecture obtained by the knapsack BU heuristics. From left to right: DNN S, DNN L, CNN S and CNN L on GOOGLE SPEECH COMMANDS. The patterns as to which computational units constitute a child subnetwork are architecture-specific.

9.2 REDS performance with 10 nested subnetworks

Table 7 and Fig. 11 show the performance of DNN, CNN and DS-CNN on GOOGLE SPEECH COMMANDS, when REDS structure comprises 10 subnetworks, compared to 4 subnetworks in the main paper. A larger number of subnetworks does not degrade model accuracy.

MACs	DNN		CNN		DS-CNN	
	Small	Large	Small	Large	Small	Large
100%	83.07 \pm 0.35	86.19 \pm 0.26	91.16 \pm 0.45	93.1 \pm 0.1	93.5 \pm 0.15	94.34 \pm 0.07
90%	82.93 \pm 0.4	86.17 \pm 0.36	90.4 \pm 0.47	92.84 \pm 0.27	93.33 \pm 0.11	94.32 \pm 0.1
80%	82.67 \pm 0.53	86.1 \pm 0.08	90.1 \pm 0.07	92.77 \pm 0.06	92.84 \pm 0.15	94.31 \pm 0.1
70%	81.67 \pm 0.53	85.78 \pm 0.11	89.43 \pm 0.41	92.25 \pm 0.47	92.64 \pm 0.24	94.21 \pm 0.14
60%	80.37 \pm 0.57	85.66 \pm 0.25	88.84 \pm 0.25	92.28 \pm 0.11	92.27 \pm 0.31	94.08 \pm 0.03
50%	78.26 \pm 0.53	85.33 \pm 0.09	88.4 \pm 0.2	92.03 \pm 0.23	91.04 \pm 0.51	93.97 \pm 0.04
40%	75.37 \pm 1.26	84.8 \pm 0.45	85.76 \pm 0.18	91.78 \pm 0.04	89.42 \pm 0.66	93.83 \pm 0.22
30%	67.76 \pm 1.59	82.66 \pm 0.18	81.92 \pm 0.49	90.52 \pm 0.54	87.63 \pm 1.03	93.59 \pm 0.14
20%	52.36 \pm 6.99	80.19 \pm 0.39	73.74 \pm 0.04	88.61 \pm 0.51	84.14 \pm 2.57	93.35 \pm 0.15
10%	23.93 \pm 5.88	50.7 \pm 7.5	58.87 \pm 5.27	81.15 \pm 1.39	58.46 \pm 3.35	90.38 \pm 0.17

Table 7: Test set accuracy [%] from Small (S) and Large (L) pretrained fully-connected (DNN), convolutional (CNN), and depth-wise separable convolutional (DS-CNN) networks taken from [53]. For each pre-trained architecture, REDS can support ten subnetworks obtained from the Knapsack BU formulation. The accuracies of the DS-CNN and CNN subnetworks do not degrade drastically until the lowest percentage of MACs considered. In contrast, the accuracies in the DNN subnetworks show a more pronounced drop from 30% MACs.

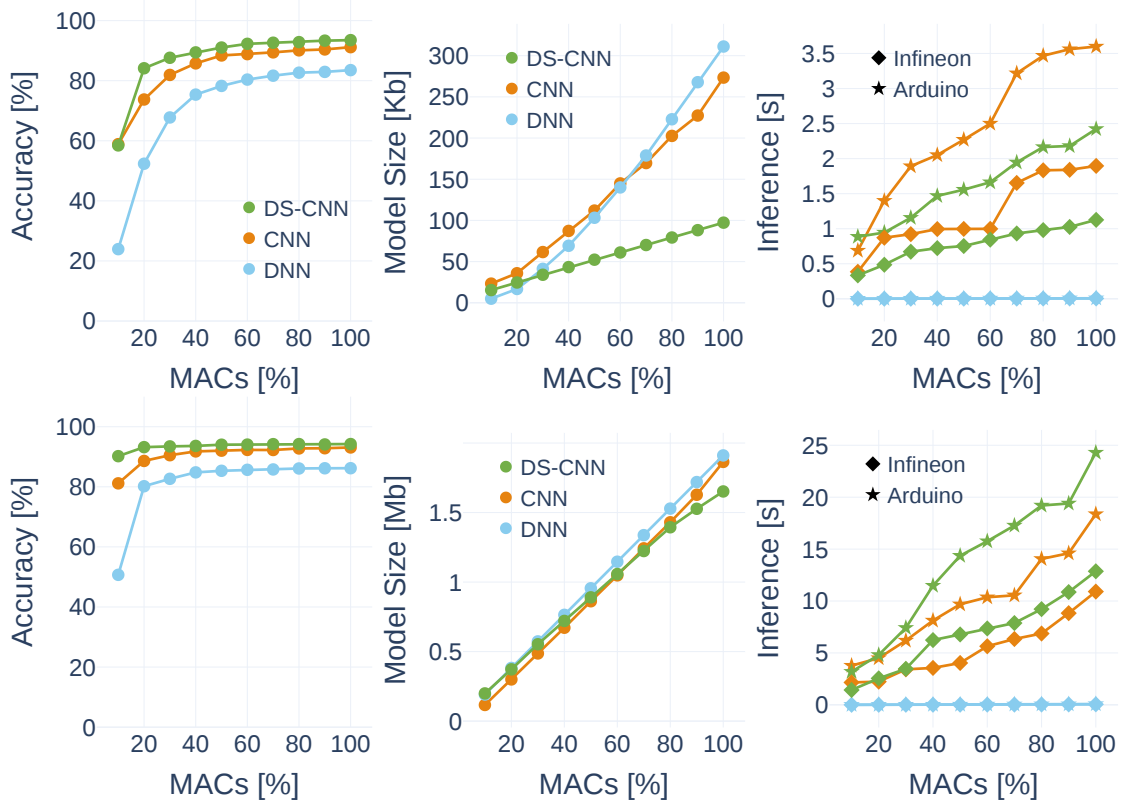


Figure 11: REDS size S (top row) and L (bottom row) architectures analysis finetuned on Google Speech Commands [50] with ten subnetworks. The plots from left to right show the subnetworks size, the subnetworks accuracy and the subnetworks inference time as a function of MAC percentage.




Figure 12: REDS fine-tuning. The parameters $\pi_{i=1}^N$ ensure the contribution of individual models to the loss aligns with the fraction of the shared weights.

9.3 REDS on FASHION-MNIST and CIFAR10

The results in Table 8 and Table 9 show REDS performance using DS-CNN architecture of size S on FASHION-MNIST and CIFAR10. BU heuristic was used to obtain the results. REDS supports a different data domain without degrading the accuracy of the pre-trained model, reported in the header row. Compared to the state-of-the-art such as μ NAS [33], REDS demonstrates a faster architecture search time for both FASHION-MNIST and CIFAR10. In the former, REDS takes 19 minutes as opposed to 3 days; in the latter, REDS takes 90 minutes as opposed to 39 days while requiring less memory for model storage for both datasets. After finding and freezing the 25% MACs subnetwork architecture, the BU heuristic takes only a few seconds to find the other 50% and 75% MACs subnetworks architectures.

MACs	Acc (%) - Pre-trained 90.59	Model Size (Kb)	Time Taken (m)
100%	91.6 \pm 0.2	128.54	–
75%	91.51 \pm 0.28	107.73	1.58 [s]
50%	90.75 \pm 0.32	87.4	5.83 [s]
25%	89.22 \pm 0.45	66.63	19 [m]

Table 8: Analysis of the BU knapsack subnetworks obtained from a depth-wise separable convolutional (DS-CNN) S network, pre-trained on FASHION-MNIST [51]. REDS supports a different data domain without degrading the accuracy of the pre-trained model, reported in the header row.

MACs	Acc (%) - Pre-trained 79.36	Model Size (Kb)	Time Taken
100%	81.07 \pm 0.71	128.54	–
75%	80.17 \pm 0.69	109.41	2.89 [s]
50%	76.72 \pm 1.37	88.01	10.59 [s]
25%	68.66 \pm 1.65	69.63	90 [m]

Table 9: The same evaluation as in Table 8 for CIFAR10 [30].